

Marketing. Which client looks like which client?

José Luis Higuera Caraveo

March 12 2022

1 Problem Statement

Marketing is essential for the growth and sustainability of any business.

Marketers can help develop the company's brand, attract customers, increase revenue and increase sales. One of the critical points for marketers is knowing their customers and identifying their needs.

By understanding the customer, marketers can launch specialized and targeted campaigns to suit each customer's specifics and needs.

By having the availability of data, referring to customer behavior, tools related to data science can be used in order to meet the specific needs of customers.

In this case study, it will be simulated that a bank in the city of New York wants to carry out a case study by launching a marketing campaign, but not just any campaign, it wants it to be specifically targeted to each type of customer. The main object is to segment them into at least 3 groups.

In the last 6 months the company has collected data regarding its customers, and these will be used for a classification model and to know this segmentation.

The data set contains the following information:

- CUSTID: Identification of the holder of the Credit Card.
- BALANCE: Amount of balance left in your account to make purchases.
- BALANCEFREQUENCY: How often the Balance is updated, score between 0 and 1 (1 = frequently updated, 0 = not frequently updated). PURCHASES: Amount of purchases made from the account.
- ONEOFFPURCHASES: Maximum purchase amount made at one time.
- INSTALLMENT PURCHASES: Amount of the purchase made in installments.
- CASHADVANCE: Cash in advance given by the user.
- FREQUENCY OF PURCHASES: How often Purchases are made, score between 0 and 1 (1 = frequent purchase, 0 = infrequent purchase).
- ONEOFFPURCHASESFREQUENCY - How often to make one-time purchases (1 = frequent purchase, 0 = infrequent purchase).
- FREQUENCY OF INSTALLMENT PURCHASES: Frequency with which installment purchases are made (1 = frequent, 0 = infrequent).
- CASHADVANCEFREQUENCY: How often the cash advance is paid.
- CASHADVANCETRX: Number of Transactions carried out with "Cash in Advanced".
- PURCHASESTRX: Number of purchase transactions made,
- CREDITLIMIT: credit card limit for the user.

- PAYMENTS: Amount of the payment made by the user.
- MINIMUM_PAYMENTS : Minimum amount of payments made by the user.
- PRCFULLPAYMENT: Percentage of the total payment paid by the user.
- TENURE: Tenure of the credit card service for the user.

The data that will be used for this study with open and public domain, are available at the following link.

[Access the data.](#)

2 Import Libraries and Data

```
[2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
matplotlib.style.use('ggplot')

from sklearn.preprocessing import StandardScaler, normalize
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
```

```
[3]: creditcard_df = pd.read_csv('./Marketing_data.csv')
creditcard_df.head()
```

```
[3]:  CUST_ID      BALANCE  BALANCE_FREQUENCY  PURCHASES  ONEOFF_PURCHASES  \
0  C10001      40.900749           0.818182      95.40           0.00
1  C10002     3202.467416           0.909091       0.00           0.00
2  C10003     2495.148862           1.000000      773.17          773.17
3  C10004     1666.670542           0.636364     1499.00         1499.00
4  C10005      817.714335           1.000000      16.00          16.00

      INSTALLMENTS_PURCHASES  CASH_ADVANCE  PURCHASES_FREQUENCY  \
0                95.4         0.000000         0.166667
1                 0.0        6442.945483         0.000000
2                 0.0         0.000000         1.000000
3                 0.0        205.788017         0.083333
4                 0.0         0.000000         0.083333

      ONEOFF_PURCHASES_FREQUENCY  PURCHASES_INSTALLMENTS_FREQUENCY  \
0                0.000000         0.083333
1                0.000000         0.000000
2                1.000000         0.000000
3                0.083333         0.000000
4                0.083333         0.000000

      CASH_ADVANCE_FREQUENCY  CASH_ADVANCE_TRX  PURCHASES_TRX  CREDIT_LIMIT  \
```

0	0.000000	0	2	1000.0
1	0.250000	4	0	7000.0
2	0.000000	0	12	7500.0
3	0.083333	1	1	7500.0
4	0.000000	0	1	1200.0

	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
0	201.802084	139.509787	0.000000	12
1	4103.032597	1072.340217	0.222222	12
2	622.066742	627.284787	0.000000	12
3	0.000000	NaN	0.000000	12
4	678.334763	244.791237	0.000000	12

We now know the columns and the data we are dealing with, we can do a quick analysis of the name of the columns, range of the data, understand the characteristics in general.

Analyzing these characteristics may not be enough to finish understanding the data, which is why a deeper analysis of it is made below.

```
[3]: creditcard_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8950 entries, 0 to 8949
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CUST_ID                               8950 non-null   object
1   BALANCE                               8950 non-null   float64
2   BALANCE_FREQUENCY                     8950 non-null   float64
3   PURCHASES                             8950 non-null   float64
4   ONEOFF_PURCHASES                      8950 non-null   float64
5   INSTALLMENTS_PURCHASES                8950 non-null   float64
6   CASH_ADVANCE                          8950 non-null   float64
7   PURCHASES_FREQUENCY                   8950 non-null   float64
8   ONEOFF_PURCHASES_FREQUENCY            8950 non-null   float64
9   PURCHASES_INSTALLMENTS_FREQUENCY      8950 non-null   float64
10  CASH_ADVANCE_FREQUENCY                 8950 non-null   float64
11  CASH_ADVANCE_TRX                       8950 non-null   int64
12  PURCHASES_TRX                         8950 non-null   int64
13  CREDIT_LIMIT                           8949 non-null   float64
14  PAYMENTS                              8950 non-null   float64
15  MINIMUM_PAYMENTS                      8637 non-null   float64
16  PRC_FULL_PAYMENT                      8950 non-null   float64
17  TENURE                                8950 non-null   int64
dtypes: float64(14), int64(3), object(1)
memory usage: 1.2+ MB
```

```
[4]: creditcard_df.describe()
```

[4]:

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	\
count	8950.000000	8950.000000	8950.000000	8950.000000	
mean	1564.474828	0.877271	1003.204834	592.437371	
std	2081.531879	0.236904	2136.634782	1659.887917	
min	0.000000	0.000000	0.000000	0.000000	
25%	128.281915	0.888889	39.635000	0.000000	
50%	873.385231	1.000000	361.280000	38.000000	
75%	2054.140036	1.000000	1110.130000	577.405000	
max	19043.138560	1.000000	49039.570000	40761.250000	

	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY	\
count	8950.000000	8950.000000	8950.000000	
mean	411.067645	978.871112	0.490351	
std	904.338115	2097.163877	0.401371	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.083333	
50%	89.000000	0.000000	0.500000	
75%	468.637500	1113.821139	0.916667	
max	22500.000000	47137.211760	1.000000	

	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	\
count	8950.000000	8950.000000	
mean	0.202458	0.364437	
std	0.298336	0.397448	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.083333	0.166667	
75%	0.300000	0.750000	
max	1.000000	1.000000	

	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	\
count	8950.000000	8950.000000	8950.000000	8949.000000	
mean	0.135144	3.248827	14.709832	4494.449450	
std	0.200121	6.824647	24.857649	3638.815725	
min	0.000000	0.000000	0.000000	50.000000	
25%	0.000000	0.000000	1.000000	1600.000000	
50%	0.000000	0.000000	7.000000	3000.000000	
75%	0.222222	4.000000	17.000000	6500.000000	
max	1.500000	123.000000	358.000000	30000.000000	

	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
count	8950.000000	8637.000000	8950.000000	8950.000000
mean	1733.143852	864.206542	0.153715	11.517318
std	2895.063757	2372.446607	0.292499	1.338331
min	0.000000	0.019163	0.000000	6.000000
25%	383.276166	169.123707	0.000000	12.000000
50%	856.901546	312.343947	0.000000	12.000000

75%	1901.134317	825.485459	0.142857	12.000000
max	50721.483360	76406.207520	1.000000	12.000000

With the last two functions applied to the dataset we can access more specific information about the data, we can analyze the mean, standard deviation, quantiles, etc., of the variables that are numerical.

We can also realize that some columns contain null data, so we have to work with them and make the decision whether to eliminate the clients that contain them or fill that data with some measure of central tendency.

Some important points are:

- Average balance is \$1564.
- Balance frequency is updated quite often, on average ~0.9.
- The average purchase is \$1000.
- The maximum non-recurring purchase amount is on average ~\$600.
- The average purchase frequency is close to 0.5.
- Average credit limit is ~\$4,500.
- Full payment percentage is 15%.
- Customers have been in the service for an average of 11 years.

```
[5]: # Analyzing the max purchase
creditcard_df[creditcard_df["ONEOFF_PURCHASES"] == 40761.25]
```

```
[5]:  CUST_ID      BALANCE  BALANCE_FREQUENCY  PURCHASES  ONEOFF_PURCHASES  \
550  C10574  11547.52001                1.0    49039.57      40761.25

      INSTALLMENTS_PURCHASES  CASH_ADVANCE  PURCHASES_FREQUENCY  \
550                8278.32    558.166886                1.0

      ONEOFF_PURCHASES_FREQUENCY  PURCHASES_INSTALLMENTS_FREQUENCY  \
550                1.0                0.916667

      CASH_ADVANCE_FREQUENCY  CASH_ADVANCE_TRX  PURCHASES_TRX  CREDIT_LIMIT  \
550                0.083333                1            101      22500.0

      PAYMENTS  MINIMUM_PAYMENTS  PRC_FULL_PAYMENT  TENURE
550  46930.59824        2974.069421            0.25      12
```

The behavior of the movements of this client justifies the amount of purchases he has made.

- The customer still has more than \$11,000 to continue shopping.
- His maximum purchase was \$40761.25, a high purchase.
- Has a high frequency of purchases (1).
- You have made 101 purchases and only 1 purchase asking for cash in advance from the bank.
- Has a credit limit of \$22,500.

In general, this client buys very frequently, is a loyal person to the bank, uses the services very frequently, so creating a marketing campaign for this type of client may not be necessary if the

objective is to encourage them to buy. One focused on using cash up front might work better.

Let's analyze the other side of the coin, that customer whose purchases are made with cash in advance.

```
[6]: creditcard_df[creditcard_df['CASH_ADVANCE'] > 47000]
```

```
[6]:
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	\
2159	C12226	10905.05381	1.0	431.93	133.5	
	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY	\		
2159	298.43	47137.21176	0.583333			
	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	\			
2159	0.25	0.5				
	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	\	
2159	1.0	123	21	19600.0		
	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE		
2159	39048.59762	5394.173671	0.0	12		

This client has the following characteristics:

- It has a high balance to make purchases.
- Its balance is constantly updated.
- Its purchases with cash in advance are greater than direct purchases.
- It has a high credit limit.


3 Data Visualization

3.1 Null and duplicate data cleaning

3.1.1 Null data

Before starting to train the Machine Learning model, it is necessary to verify if the data has null data. To do this, the following visualization is performed to verify it.

```
[7]: sns.heatmap(creditcard_df.isnull(), yticklabels=False, cbar=False, cmap='Blues')  
plt.show()
```



The visualization helps us understand that the `MINIMUM_PAYMENTS` column has missing data. Another way to check it is as follows.

```
[8]: creditcard_df.isnull().sum()
```

```
[8]: CUST_ID          0
      BALANCE         0
      BALANCE_FREQUENCY  0
      PURCHASES       0
      ONEOFF_PURCHASES  0
      INSTALLMENTS_PURCHASES  0
      CASH_ADVANCE    0
```

PURCHASES_FREQUENCY	0
ONEOFF_PURCHASES_FREQUENCY	0
PURCHASES_INSTALLMENTS_FREQUENCY	0
CASH_ADVANCE_FREQUENCY	0
CASH_ADVANCE_TRX	0
PURCHASES_TRX	0
CREDIT_LIMIT	1
PAYMENTS	0
MINIMUM_PAYMENTS	313
PRC_FULL_PAYMENT	0
TENURE	0

dtype: int64

It is important to carry out the previous operation since it gives us the most detailed information of the missing data. In the visualization it is not possible to appreciate that the CREDIT_LIMIT column also has missing data, only 1, which also has to be worked on.

For this case study, the missing data will be filled in with the mean of the corresponding column.

```
[6]: # Removing the missing ones for the MINIMUM PAYMENTS column
creditcard_df.loc[(creditcard_df['MINIMUM_PAYMENTS'].isnull() == True),
↳ 'MINIMUM_PAYMENTS'] = creditcard_df['MINIMUM_PAYMENTS'].mean()

# Removing the missing ones for the CREDIT_LIMIT column
creditcard_df.loc[(creditcard_df['CREDIT_LIMIT'].isnull() == True),
↳ 'CREDIT_LIMIT'] = creditcard_df['CREDIT_LIMIT'].mean()

# Checking for missing data
creditcard_df.isnull().sum()
```

[6]: CUST_ID	0
BALANCE	0
BALANCE_FREQUENCY	0
PURCHASES	0
ONEOFF_PURCHASES	0
INSTALLMENTS_PURCHASES	0
CASH_ADVANCE	0
PURCHASES_FREQUENCY	0
ONEOFF_PURCHASES_FREQUENCY	0
PURCHASES_INSTALLMENTS_FREQUENCY	0
CASH_ADVANCE_FREQUENCY	0
CASH_ADVANCE_TRX	0
PURCHASES_TRX	0
CREDIT_LIMIT	0
PAYMENTS	0
MINIMUM_PAYMENTS	0
PRC_FULL_PAYMENT	0
TENURE	0

dtype: int64

3.1.2 Duplicated data

Another important process when analyzing the data is if we have duplicate data.

```
[4]: duplicates = creditcard_df.duplicated().sum()

print('Duplicated data: {}'.format(duplicates))
```

Duplicated data: 0

In this case we do not have duplicate data.

3.2 Cleaning of unnecessary data

The CUST_ID column is not necessary for modeling, so we proceed to eliminate it, in order not to continue working with it.

```
[7]: creditcard_df.drop('CUST_ID', axis=1, inplace=True)
creditcard_df.head()
```

```
[7]:
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	\
0	40.900749	0.818182	95.40	0.00	
1	3202.467416	0.909091	0.00	0.00	
2	2495.148862	1.000000	773.17	773.17	
3	1666.670542	0.636364	1499.00	1499.00	
4	817.714335	1.000000	16.00	16.00	

	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY	\
0	95.4	0.000000	0.166667	
1	0.0	6442.945483	0.000000	
2	0.0	0.000000	1.000000	
3	0.0	205.788017	0.083333	
4	0.0	0.000000	0.083333	

	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	\
0	0.000000	0.083333	
1	0.000000	0.000000	
2	1.000000	0.000000	
3	0.083333	0.000000	
4	0.083333	0.000000	

	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	\
0	0.000000	0	2	1000.0	
1	0.250000	4	0	7000.0	
2	0.000000	0	12	7500.0	
3	0.083333	1	1	7500.0	

4		0.000000	0	1	1200.0
	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE	
0	201.802084	139.509787	0.000000	12	
1	4103.032597	1072.340217	0.222222	12	
2	622.066742	627.284787	0.000000	12	
3	0.000000	864.206542	0.000000	12	
4	678.334763	244.791237	0.000000	12	

Removing the CUST_ID column leaves us with 17 to work on.

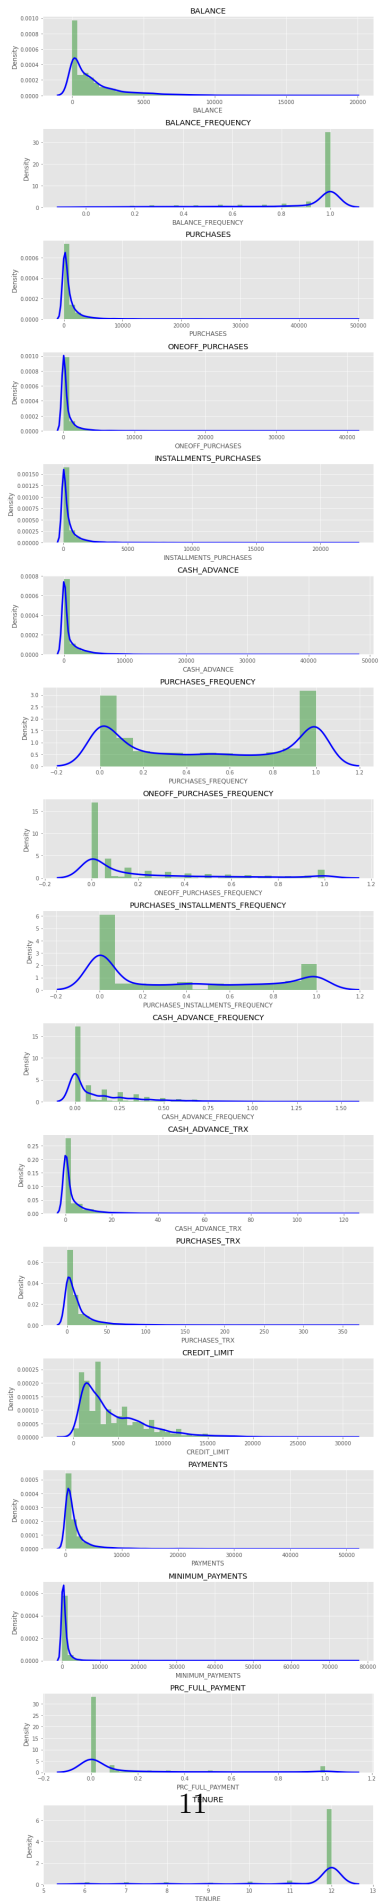
3.3 Data Visualization

3.3.1 Histograms and KDE Density

A useful way to look at the density of a variable is with a KDE chart, which helps us visualize the probability density of a continuous variable.

```
[13]: plt.figure(figsize = (10, 50))
      for i in range(len(creditcard_df.columns)):
          plt.subplot(len(creditcard_df.columns), 1, i+1)
          sns.distplot(creditcard_df[creditcard_df.columns[i]], kde_kws = {"color": "b", "lw": 3, "label": "KDE"}, hist_kws={"color": "g"})
          plt.title(creditcard_df.columns[i])

      plt.tight_layout()
      plt.show()
```




- Average balance is \$1,500.
- 'Balance_Frequency' for many users is updated very frequently ~ 1 .
- For the 'PURCHASES_FREQUENCY' field, there are two different groups of customers.
- For the 'ONEOFF_PURCHASES_FREQUENCY' and 'PURCHASES_INSTALLMENT_FREQUENCY' fields, the vast majority of users do not pay everything at once or in installments.
- Very few customers pay their debt in full 'PRC_FULL_PAYMENT' ~ 0 .
- The average credit limit is around \$4,500.
- Most customers have been using the service for ~ 11 years.

3.3.2 Correlation matrix

We need to understand how the variables are correlated, this helps us find possible trends in the data.

```
[14]: correlations = creditcard_df.corr()

f, ax, = plt.subplots(figsize=(20,20))
sns.heatmap(correlations, annot=True)
plt.show()
```



4 Training Model

4.1 K-Means

Before starting to model, it is necessary to scale the data, this to prevent variables with a larger range from dominating versus others with smaller domains.

```
[8]: scaler = StandardScaler()
creditcard_df_scaled = scaler.fit_transform(creditcard_df)
```


4.1.1 Optimal number of centroids (Elbow method).

When applying a K-Means model, it is necessary to find the optimal number of clusters in which the data is going to be divided, there is no mathematical formula to process this data, so far, the best way to do it is with the visualization of the elbow method. This method takes as a scoop the intra-cluster variance between the centroid and the data that compose it. At a large number of k (centroids), less variance. The elbow method, visualizing this variance, takes as optimal k the one whose difference between k and k + 1 no longer improves considerably.

```
[9]: # Find optimal K
scores_1 = []
range_values = range(1,20)

for i in range_values:
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(creditcard_df_scaled)
    scores_1.append(kmeans.inertia_)

# Visualize to choose the optimal K
plt.plot(range_values, scores_1, 'kx-')
plt.title('Elbow Method')
plt.xlabel('Clusters')
plt.ylabel('WCSS(k)')
plt.show()
```



With the graph we can see that in 4 clusters is where the elbow of the curve is formed. However, the values do not reduce to a linear form until the 8th cluster. Let us choose a number of clusters equal to 8.

4.1.2 Training

```
[17]: k = 8

kmeans = KMeans(n_clusters=k)
kmeans.fit(creditcard_df_scaled)
labels = kmeans.labels_
```

Cluster centers are those centroids that help separate data into different segments. To visualize them we include them in a Dataframe.

```
[18]: cluster_centers = pd.DataFrame(data = kmeans.cluster_centers_, columns = creditcard_df.columns)
cluster_centers
```

```
[18]:
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	\
0	0.901817	0.466986	2.270963	1.756936	
1	0.019493	0.403153	-0.361863	-0.246971	
2	1.698325	0.393098	-0.215463	-0.154529	
3	1.923051	0.337717	11.212042	10.600367	

4	-0.701229	-2.144116	-0.311099	-0.235720
5	-0.336050	-0.347078	-0.289267	-0.215966
6	-0.165253	0.392196	0.453349	0.593167
7	-0.364778	0.333613	-0.037381	-0.244339

	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY	\
0	2.141920	-0.195512	1.158629	
1	-0.401779	-0.086621	-0.867503	
2	-0.225632	2.025668	-0.471452	
3	7.033118	0.419625	1.046983	
4	-0.302414	-0.321905	-0.556586	
5	-0.286835	0.068284	-0.203078	
6	-0.017967	-0.333914	0.943302	
7	0.360316	-0.363589	0.990669	

	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	\
0	1.583889	1.226198	
1	-0.410513	-0.758672	
2	-0.210500	-0.409161	
3	1.915501	0.981334	
4	-0.444989	-0.439730	
5	-0.288661	-0.224549	
6	1.878357	0.089014	
7	-0.387079	1.206081	

	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	\
0	-0.312181	-0.212097	2.781452	1.238947	
1	0.115631	-0.020700	-0.486861	-0.305126	
2	1.920837	1.941432	-0.263115	1.040171	
3	-0.258912	0.061229	5.362438	3.044064	
4	-0.520844	-0.376103	-0.419790	-0.177161	
5	0.308663	0.000996	-0.388117	-0.567159	
6	-0.407665	-0.323378	0.523732	0.373578	
7	-0.475238	-0.361153	0.187666	-0.260925	

	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
0	1.290295	0.441655	0.304778	0.334182
1	-0.248169	-0.008412	-0.456474	0.271801
2	0.828342	0.557352	-0.392330	0.071341
3	8.098975	1.120318	1.110132	0.310863
4	-0.202048	-0.256658	0.281550	0.199199
5	-0.392680	-0.209145	0.014011	-3.203733
6	0.086557	-0.162605	0.406347	0.261047
7	-0.216886	-0.032660	0.313849	0.257637

The above data is scaled, so it's hard to really understand what it means. For this, we will apply the inverse transformation of the scaling to obtain the real values and thus be able to analyze more

clearly and adequately.

```
[19]: cluster_centers = scaler.inverse_transform(cluster_centers)
cluster_centers = pd.DataFrame(data = cluster_centers, columns = [creditcard_df.
→columns])
cluster_centers
```

```
[19]:      BALANCE BALANCE_FREQUENCY PURCHASES ONEOFF_PURCHASES \
0  3441.530986      0.987896  5855.151608      3508.591111
1  1605.047605      0.972774   230.077907      182.515426
2  5099.393953      0.970392   542.864477      335.950907
3  5567.142164      0.957273 24957.905000     18186.875667
4   104.925267      0.369349   338.537483      201.190254
5   865.015978      0.795051   385.181720      233.977974
6  1220.514994      0.970178  1971.792676     1576.972447
7   805.220083      0.956301   923.338824      186.885283
```

```
      INSTALLMENTS_PURCHASES CASH_ADVANCE PURCHASES_FREQUENCY \
0          2347.978936      568.874079      0.955365
1           47.744156      797.223294      0.142179
2          207.031791     5226.790667      0.301134
3         6771.029333     1858.844605      0.910556
4          137.598754      303.821813      0.266966
5          151.686061     1122.064941      0.408846
6          394.820228      278.637458      0.868943
7          736.896637      216.408238      0.887954
```

```
      ONEOFF_PURCHASES_FREQUENCY PURCHASES_INSTALLMENTS_FREQUENCY \
0              0.674962              0.851760
1              0.079994              0.062922
2              0.139661              0.201826
3              0.773889              0.754444
4              0.069709              0.189677
5              0.116344              0.275196
6              0.762808              0.399814
7              0.086985              0.843765
```

```
      CASH_ADVANCE_FREQUENCY CASH_ADVANCE_TRX PURCHASES_TRX CREDIT_LIMIT \
0              0.072674          1.801418      83.846336      9002.245863
1              0.158283          3.107562      2.608297      3384.275575
2              0.519523         16.497674      8.169767      8279.016913
3              0.083333          3.666667     148.000000     15570.000000
4              0.030918          0.682203      4.275424      3849.863936
5              0.196911          3.255627      5.062701      2430.891398
6              0.053566          1.042009      27.727854      5853.677875
7              0.040044          0.784226      19.374504      3545.099307
```

	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
0	5468.421612	1893.464949	0.242857	11.964539
1	1014.718178	844.603245	0.020204	11.881057
2	4131.114001	2163.092995	0.038965	11.612791
3	25178.882690	3475.059479	0.478409	11.933333
4	1148.234177	266.075424	0.236063	11.783898
5	596.373827	376.802926	0.157813	7.229904
6	1983.717894	485.262318	0.272564	11.866667
7	1105.280930	788.094852	0.245510	11.862103

Let's analyze the most relevant clusters.

- Second Cluster of Customers: These are customers who use the service very little, have a purchase frequency of just 0.14, although they have a balance of \$1,605, so they are considered as those customers who save.
- Fourth Cluster of Clients: They have a balance greater than the average, but without being very high, high purchase frequency, their purchases in a single transaction are high, and they also use credit frequently.
- Sixth cluster of Clients: They are those whose level of purchase is the highest, their maximum purchase amount at one time is the highest. That is, those customers with a greater flow of purchases. Therefore, those who pay the most taxes.
- Seventh cluster of Clients: They are those who contribute more cash, that is, they do not ask for it in advance from the bank, so their transactions with Cash in Advance are low.
- Eighth cluster of Clients: These are clients that are similar to the previous ones, only that they buy less and their account balance is not very high.

Now we can add to the original dataset to which cluster each client corresponds.

```
[20]: creditcard_df_cluster = pd.concat([creditcard_df, pd.DataFrame({'cluster':  
↪labels})], axis = 1)  
creditcard_df_cluster.head()
```

```
[20]:
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	\
0	40.900749	0.818182	95.40	0.00	
1	3202.467416	0.909091	0.00	0.00	
2	2495.148862	1.000000	773.17	773.17	
3	1666.670542	0.636364	1499.00	1499.00	
4	817.714335	1.000000	16.00	16.00	

	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY	\
0	95.4	0.000000	0.166667	
1	0.0	6442.945483	0.000000	
2	0.0	0.000000	1.000000	
3	0.0	205.788017	0.083333	
4	0.0	0.000000	0.083333	

	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	\
0	0.000000	0.083333	

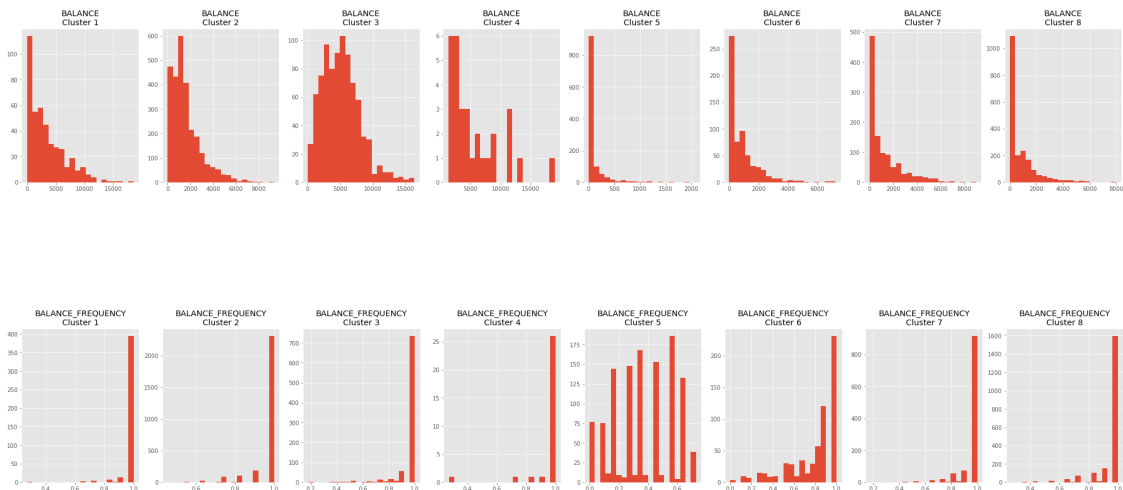
1	0.000000	0.000000
2	1.000000	0.000000
3	0.083333	0.000000
4	0.083333	0.000000






	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT \
0	0.000000	0	2	1000.0
1	0.250000	4	0	7000.0
2	0.000000	0	12	7500.0
3	0.083333	1	1	7500.0
4	0.000000	0	1	1200.0

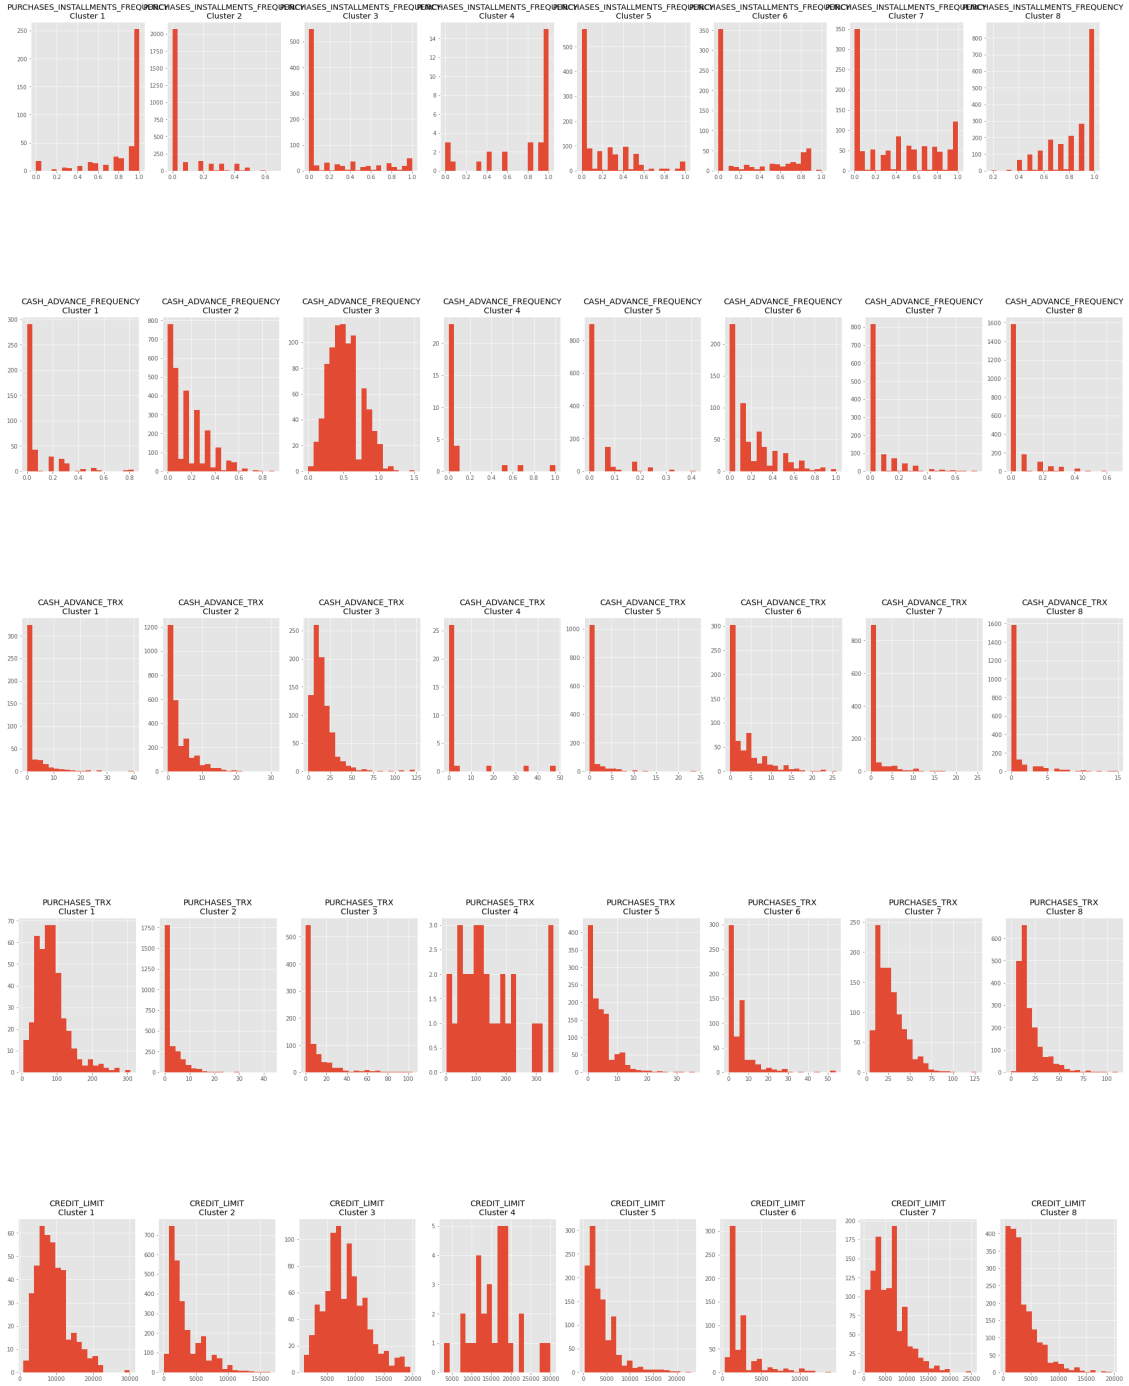
	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE	cluster
0	201.802084	139.509787	0.000000	12	1
1	4103.032597	1072.340217	0.222222	12	2
2	622.066742	627.284787	0.000000	12	6
3	0.000000	864.206542	0.000000	12	1
4	678.334763	244.791237	0.000000	12	1


Visualizing the histogram of each variable with respect to the clusters.

```
[22]: for i in creditcard_df.columns:
plt.figure(figsize=(35,5))
for j in range(8):
plt.subplot(1,8,j+1)
cluster = creditcard_df_cluster[creditcard_df_cluster['cluster'] == j]
cluster[i].hist(bins = 20)
plt.title('{}\nCluster {}'.format(i, j+1))
plt.show()
```









In the previous visualization we can see how each variable behaves with respect to the cluster, as well as the frequencies assigned to each cluster. for example:

- Customers with the highest balance are classified in clusters 3 and 4.
- Customers who make high purchases at a single display are in cluster 4.
- Customers who update their balance the least are in cluster 5.
- Customers in clusters 3 and 6 are the ones who ask for more cash in advance from the bank.

4.2 Principal Component Analysis

Analyzing a data set where many variables are involved, in this case study 17, can become complicated.

The PCA or ACP in Spanish, helps us to reduce the dimensionality of the problem, this by reducing the number of variables, such that the least possible variance of the data is lost and thus the least possible information is lost.

The objective of this case study is to reduce it to 2 main components, this in order to be able to visualize the users in a scatter graph and to make it easier to analyze it.

```
[23]: pca = PCA(n_components=2)
principal_comp = pca.fit_transform(creditcard_df_scaled)
```

```
[24]: pca_df = pd.DataFrame(data=principal_comp, columns=['pca_1', 'pca_2'])
pca_df.head()
```

```
[24]:      pca_1      pca_2
0 -1.682222 -1.076444
1 -1.138299  2.506500
2  0.969687 -0.383521
3 -0.873628  0.043176
4 -1.599436 -0.688578
```

We already have the data projected to only 2 dimensions, now we will concatenate the clusters to which each client belongs.


```
[25]: pca_df = pd.concat([pca_df, pd.DataFrame({'cluster': labels})], axis = 1)
pca_df.head()
```

```
[25]:      pca_1      pca_2  cluster
0 -1.682222 -1.076444         1
1 -1.138299  2.506500         2
2  0.969687 -0.383521         6
3 -0.873628  0.043176         1
4 -1.599436 -0.688578         1
```

With a dataframe of only two features we can visualize the result of the K-Means algorithm more clearly.

```
[26]: plt.figure(figsize=(10,10))
ax = sns.scatterplot(x='pca_1', y='pca_2', hue='cluster', data=pca_df,
                    palette=['r', 'g', 'b', 'pink', 'yellow', 'gray', 'purple', 'k'])

plt.show()
```



As we can see, we can now visualize how the K-Means algorithm has clustered the data.

The next step is to put this information into practice and update the number k of clusters with respect to the behavior of the market.

The K-Means algorithm, being unsupervised, is difficult to predict with high accuracy the correct number of clusters, the elbow method is useful and helps to give a first impression, but it will be the behavior of the algorithm in practice that help maintain the code and make the necessary changes.

4.3 Autoencoders

Analyzing a large number of variables to understand what is being studied can be a difficult task.

We will use the autoencoders to reduce the dimensionality of the problem and thus make it easy to analyze the variables.


```
[27]: from tensorflow.keras.layers import Input, Add, Dense, Activation, \
      ↪ZeroPadding2D, BatchNormalization, Flatten, Conv2D, AveragePooling2D, \
      ↪MaxPooling2D, Dropout
      from tensorflow.keras.models import Model, load_model
      from tensorflow.keras.initializers import glorot_uniform
      from tensorflow.keras.optimizers import SGD
```

4.3.1 Autoencoder Architecture

For this case study we will start from 17 variables (The 17 original columns of the dataset) and we will compress them to obtain 10.

```
[32]: input_df = Input(shape = (17,))
      encoding_dim = 7

      # We make a first reduction to 7 variables
      x = Dense(encoding_dim, activation='relu')(input_df)
      # We work with that reduction
      x = Dense(500, activation='relu', kernel_initializer='glorot_uniform')(x)
      x = Dense(500, activation='relu', kernel_initializer='glorot_uniform')(x)
      x = Dense(2000, activation='relu', kernel_initializer='glorot_uniform')(x)

      # We increase to 10 variables
      encoded = Dense(10, activation='relu', kernel_initializer='glorot_uniform')(x)

      # We work with those 10 variables in an inverse way, biasing to not have a
      ↪layer of 500 neurons
      x = Dense(2000, activation='relu', kernel_initializer='glorot_uniform')(encoded)
      x = Dense(500, activation='relu', kernel_initializer='glorot_uniform')(x)

      # We decode to obtain again the 17 variables
      decoded = Dense(17, kernel_initializer='glorot_uniform')(x)

      # Model to encode and decode
      autoencoder = Model(input_df, decoded)
      # Model only for encoding
      encoder = Model(input_df, encoded)

      autoencoder.compile(optimizer='adam', loss='mean_squared_error')
```

```
[33]: # A summary of how the model compiles the information
      autoencoder.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #
=====		

input_2 (InputLayer)	[(None, 17)]	0

dense_8 (Dense)	(None, 7)	126

dense_9 (Dense)	(None, 500)	4000

dense_10 (Dense)	(None, 500)	250500

dense_11 (Dense)	(None, 2000)	1002000

dense_12 (Dense)	(None, 10)	20010

dense_13 (Dense)	(None, 2000)	22000

dense_14 (Dense)	(None, 500)	1000500

dense_15 (Dense)	(None, 17)	8517
=====		
Total params: 2,307,653		
Trainable params: 2,307,653		
Non-trainable params: 0		

```
[41]: # Training
hist = autoencoder.fit(creditcard_df_scaled, creditcard_df_scaled,
    ↪ batch_size=64, epochs=15,
        verbose=1)
```

```
Epoch 1/15
140/140 [=====] - 10s 69ms/step - loss: 0.0588
Epoch 2/15
140/140 [=====] - 13s 93ms/step - loss: 0.0566
Epoch 3/15
140/140 [=====] - 8s 60ms/step - loss: 0.0501
Epoch 4/15
140/140 [=====] - 8s 60ms/step - loss: 0.0466
Epoch 5/15
140/140 [=====] - 6s 46ms/step - loss: 0.0447
Epoch 6/15
140/140 [=====] - 14s 102ms/step - loss: 0.0635
Epoch 7/15
140/140 [=====] - 13s 90ms/step - loss: 0.0523
Epoch 8/15
140/140 [=====] - 10s 69ms/step - loss: 0.0414
Epoch 9/15
140/140 [=====] - 13s 92ms/step - loss: 0.0369
Epoch 10/15
140/140 [=====] - 12s 88ms/step - loss: 0.0336
```

```
Epoch 11/15
140/140 [=====] - 8s 59ms/step - loss: 0.0342
Epoch 12/15
140/140 [=====] - 12s 87ms/step - loss: 0.0402
Epoch 13/15
140/140 [=====] - 10s 75ms/step - loss: 0.0462
Epoch 14/15
140/140 [=====] - 11s 78ms/step - loss: 0.0316
Epoch 15/15
140/140 [=====] - 12s 84ms/step - loss: 0.0293
```

```
[42]: # Save weights of the trained autoencoder
autoencoder.save_weights('autoencoder.h5')
```

Now that we have the autoencoder we train, it is time to reduce the variables from 17 to 10.

We will call this new data group pred.


```
[43]: pred = encoder.predict(creditcard_df_scaled)
```

We now return to the K-Means algorithm, this to now apply it to the reduction of variables.

```
[44]: # Find optimal K
scores_1 = []
range_values = range(1,20)

for i in range_values:
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(pred)
    scores_1.append(kmeans.inertia_)

# Visualize to choose the optimal K
plt.plot(range_values, scores_1, 'kx-')
plt.title('Elbow Method')
plt.xlabel('Clusters')
plt.ylabel('WCSS(k)')
plt.show()
```



Now that we have the autoencoder we train, it is time to reduce the variables from 17 to 10.

We will call this new data group pred.

```
[45]: kmeans = KMeans(5)
      kmeans.fit(pred)
      labels = kmeans.labels_
      y_kmeans = kmeans.fit_predict(pred)

      df_cluster_ae = pd.concat([creditcard_df, pd.DataFrame({'cluster': labels})],
      ↪axis=1)
      df_cluster_ae.head()
```

```
[45]:
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	\
0	40.900749	0.818182	95.40	0.00	
1	3202.467416	0.909091	0.00	0.00	
2	2495.148862	1.000000	773.17	773.17	
3	1666.670542	0.636364	1499.00	1499.00	
4	817.714335	1.000000	16.00	16.00	

	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY	\
0	95.4	0.000000	0.166667	
1	0.0	6442.945483	0.000000	
2	0.0	0.000000	1.000000	

3	0.0	205.788017	0.083333
4	0.0	0.000000	0.083333

	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	\
0	0.000000	0.083333	
1	0.000000	0.000000	
2	1.000000	0.000000	
3	0.083333	0.000000	
4	0.083333	0.000000	

	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	\
0	0.000000	0	2	1000.0	
1	0.250000	4	0	7000.0	
2	0.000000	0	12	7500.0	
3	0.083333	1	1	7500.0	
4	0.000000	0	1	1200.0	

	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE	cluster
0	201.802084	139.509787	0.000000	12	4
1	4103.032597	1072.340217	0.222222	12	2
2	622.066742	627.284787	0.000000	12	4
3	0.000000	864.206542	0.000000	12	1
4	678.334763	244.791237	0.000000	12	4

To visualize these new clusters we apply PCA again to 2 dimensions.


```
[46]: pca = PCA(n_components=2)
pca_comp2 = pca.fit_transform(pred)
pca_df = pd.DataFrame(data=pca_comp2, columns=['pca_1', 'pca_2'])

# We concatenate the calculated clusters with k-means
pca_df = pd.concat([pca_df, pd.DataFrame({'cluster': labels})], axis=1)
pca_df.head()
```

```
[46]:      pca_1      pca_2  cluster
0 -1.776384 -0.187245         4
1  0.192772  1.718288         2
2 -0.488257 -1.082752         4
3 -0.229187 -0.191664         1
4 -2.063338 -0.241279         4
```

```
[47]: plt.figure(figsize=(10,10))
ax = sns.scatterplot(x='pca_1', y='pca_2', hue='cluster', data=pca_df,
                    palette=['r', 'g', 'b', 'purple', 'k'])

plt.show()
```



The visualization makes it clear to us how the data is segmented with 5 clusters.

The next steps are to deploy the algorithm in production, observe the behavior and improve it accordingly. Remember that there is no exact formula to calculate the optimal clusters, so what remains is to improve as customers respond to different campaigns.